

Using PCA to Efficiently Represent State Spaces

Editor: Editor

Abstract

Reinforcement learning algorithms must cope with the exponential growth of states and actions when learning control policies in high-dimensional spaces, i.e., the curse of dimensionality. By projecting the agent’s state onto a low-dimensional manifold, we can represent the state space in a smaller and more efficient representation. By using this representation during learning, the agent can converge to a good policy much faster than if it used the full representation. We test this idea in the Mario Benchmarking Domain. When using dimensionality reduction in Mario, learning converges much faster to a good policy. But, there is a critical convergence-performance trade-off — by projecting onto a low-dimensional manifold, important data is ignored. In this paper, we explore this trade-off of convergence and performance. We find that learning in as few as 4 dimensions (instead of 9), we can significantly improve the learning performance relative to using the full dimensional space.

1. Introduction

Learning in high dimensional spaces is necessary and difficult in robotic applications. State and action spaces in robotics become large, continuous, and scale exponentially with the number of joints. This leads to the *curse of dimensionality*.

One approach to tackling this problem is through the use of transfer learning (Taylor and Stone, 2009; Pan and Yang, 2010) and learning from demonstration (Argall et al., 2009) approaches. Transfer learning reduces computational complexity by learning in a simple domain, and transferring that knowledge to a more complex domain. In transfer learning there are three key research questions: What to transfer, how to transfer, and when to transfer. Each of these questions are difficult to answer and all are domain dependent. When the source domain and the target domain are loosely related, straight forward transfer learning does not work, and can lead to worse performance (Taylor et al., 2007; Pan and Yang, 2010). Transfer learning also requires a computable mapping from the source domain to the target domain, which is not always possible.

Learning from demonstration (LfD) methods speed up convergence time by bootstrapping learning with demonstrations. LfD learns a policy using examples or demonstrations provided by a human or a robotic teacher. This method extracts state-action pairs from these demonstrations to bootstrap learning. However, these demonstrations must be consistent and accurately represent solving the task. These methods also solve a specific complex task, rather than solve for general control (Argall et al., 2009).

In this work, we focus on the core problem of dealing with large state spaces. We approach this issue by projecting the full state space onto a low-dimensional manifold. We use Principal Component Analysis to find this transform, and use it during each learning iteration. In this way, we perform learning only in the low-dimensional space. This leads to a critical trade-off: by projecting

onto a low-dimensional manifold, we are throwing out low variance, yet potentially important data. However, learning in a low dimensional space can converge to a good, although possibly suboptimal, policy much faster than when learning in the original, high-dimensional space. In this paper, we explore this trade-off of convergence and performance.

We organize the rest of this paper as follows. Section 2 describes related work in dimensionality reduction in machine learning and reinforcement learning. In Section 3 we describe our work with PCA, dimensionality reduction and reinforcement learning. Section 4 introduces the Mario Benchmarking Domain and our learning approach. Experimental results are in Section 5, followed by conclusion and related work in Sections 6 and 7.

2. Related Work

Previous work in dimensionality reduction focuses on reducing the space for classification or function approximation. PCA is effective in many machine learning and data mining applications to extract features from large data sets (Pechenizkiy et al., 2003; Turk and Pentland, 1991). Rather than using PCA for feature extraction in large data sets, we use PCA to reduce the dimensionality of the state space during learning.

Swinehart and Abbott (2005) used a similar approach for function approximation with neural networks. They find that they can reduce convergence time for reinforcement-based, random walk learning by reducing the dimension of the parameter space. Liu and Mahadevan (2011) also use dimensionality reduction to compute policies in a low dimensional subspace. Liu computes the low dimensional subspace from a high-dimensional space through random projections. They also reduce convergence time in continuous state spaces.

The work of Ahmadi et al. (2007) is similar in spirit because the agent learns using a subset of the state variables (i.e., the agent initially ignores many of the state variables). Although this work focuses on learning in discrete spaces, it could in principle also be done in continuous spaces (e.g., Taylor and Stone (2007)). However, this process is hand-coded. In IFSA, a human must specify 1) which state variables to begin learning on, 2) when the agent should begin considering a new state variable, and 3) which state variable should be added to the representation. Given this information, the agent is able to learn in a lower dimensional space and gradually complexify the representation, using transfer learning to bootstrap learning in each new state space.

3. Learning in High-Dimensional Spaces

Robots must perform complex manipulation tasks successfully to be useful in many real-world scenarios. These complex manipulations need high degree-of-freedom arms and manipulators. For example, the PR2 robot has two 7 DoF arms. When learning position, velocity, and acceleration control, this leads to a 21 dimensional state space per arm. Learning in these high-dimensional spaces is computationally intractable without optimization techniques.

To learn in high-dimensional state spaces, our algorithm first computes a transform between the high-dimensional state space and a lower-dimensional space (Figure 1). To perform this computation, we need trajectories across a representative set of the agent’s state space. We can then use any dimensionality reduction technique to learn the transform. In this work, we use Principal Component Analysis (PCA) (Shlens, 2005).

PCA identifies patterns in data and reduces the dimensions of the dataset with minimal loss of information. It does this by computing a transform to convert correlated data to linearly uncorrelated data. This transformation ensures that the first principal component has the largest possible variance. Each additional component has the largest possible variance, uncorrelated with all previous components. Essentially, PCA represents as much of the demonstrated state space as possible in a lower dimension. This transform is given by:

$$T = XW \quad (1)$$

where X is the demonstrated data, W is a p by p matrix whose columns are eigenvectors of $X^T X$ and p is the number of principle components (in this case, the number of dimensions).

To transform to any arbitrary dimension, k , we can choose k eigenvectors from W with the largest eigenvalues to form a p by k dimensional matrix W_k :

$$T_k = XW_k \quad (2)$$

We then use reinforcement learning to learn the policy in the new manifold. All learning is in the lower-dimensional space. For each learning step, we project state x down to the lower-dimensional space k :

$$x_k = W_k^T x \quad (3)$$

We can now compute the action using the chosen learning algorithm, and execute that action in the simulation. The simulation calculates the new state given the executed action, and we can project that state down to the same lower-dimensional space. We can then perform a learning update (Figure 1).

By learning in a smaller space, reinforcement learning algorithms will converge much faster. However, PCA cannot represent all the variance in all the demonstrations. Therefore, given an infinite number of learning iterations, the converged learning performance in the full space will always be better than or at least as good as in the smaller space. This leads to a critical trade-off. By projecting onto a low-dimensional manifold, we are throwing out low variance, yet possibly important data. In practice, we find that learning can still converge to a good, but suboptimal, policy much faster than when learning in the full state space.

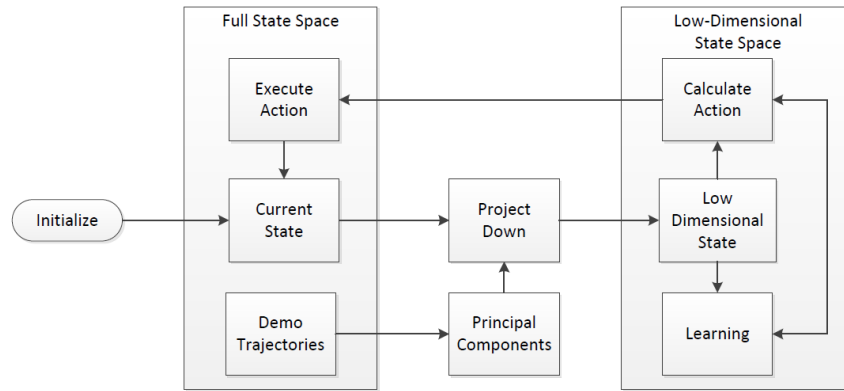


Figure 1: PCA+RL Flowchart

4. Mario Benchmark Problem

The Mario benchmark problem (Karakovskiy and Togelius, 2012) is based on Infinite Mario Bros, which is a public reimplementation of the original 80’s game Super Mario Bros[®]. In this task, Mario needs to collect as many points as possible, this is done by killing an enemy (10), devouring a mushroom (58) or a fireflower (64), grabbing a coin (16), finding a hidden block (24), finishing the level (1024), getting hurt by a creature (−42) or dying (−512). The actions available to Mario correspond to the buttons on the NES controller, which are (left, right, no direction), (jump, do not jump), and (run/fire, do not run/fire). Mario can take one action from each of these groups simultaneously, resulting in 12 distinct combined or ‘super’ actions. The state space in Mario is quite complex, as Mario observes the exact locations of all enemies on the screen and their type. He also observes all information about himself, such as what mode he is in (small, big, fire). Lastly, he has a gridlike receptive field in which each cell indicates what type of object is in it (such as a brick, a coin, a mushroom, a goomba (enemy), etc.). A screenshot is shown in Figure 2.

Our reinforcement learning agent for Mario is inspired by previous work in the domain (Liao et al., 2012; Brys et al., 2014). We use a $Q(\lambda)$ -learner with a tabular state representation and $\alpha=0.01$, $\lambda=0.5$, $\gamma=0.9$, and $\epsilon=0.05$. The part of the state space the agent considers consists of the following state variables:

- is Mario able to jump (0, 1)
- is Mario on the ground (0, 1)
- is Mario able to shoot fireballs (0, 1)
- Mario’s current direction, 8 directions and standing still (0 – 8)
- enemies close by (within one gridcell) in 8 directions ($2^8 \rightarrow 0-255$)
- enemies at midrange (within one to three gridcells) in 8 directions ($2^8 \rightarrow 0 - 255$)
- whether there is an obstacle in four vertical gridcells in front of Mario ($2^4 \rightarrow 0 - 16$)
- closest enemy position within 21x21 grid surrounding Mario + 1 for absent enemy ($21^2 + 1 \rightarrow 0 - 441$)

This makes for a total of 3.24×10^{10} possible states, and 4×10^{11} Q -values (12 possible actions per state). The size of the state space is not a problem computationally, because this state-space is sparsely visited, with the majority of the state-visits to a small set of states (Liao et al., 2012). To further improve learning, previous work has investigated using demonstrations in the Mario domain to speed up reinforcement learning, via shaping the reward using demonstration (Brys et al., 2015) and learning a reward function using inverse reinforcement learning (Lee et al., 2014).

5. Results

In order to evaluate our approach, we run every learning episode on a procedurally generated level based on a random seed $\in [0, 10^6]$, with difficulty 0. Mario randomly starts as one of {small, large, fire} for each episode. Making an agent learn to play Mario this way helps avoid overfitting on a



Figure 2: A screenshot of Mario.

specific level, and makes for a more generally applicable Mario agent. All results are averaged over 100 different trials.

We found the principal components using data we gathered from a single demonstration. There were 2905 sample points in this demonstration. We then scaled all values between 0 and 1 to remove scaling variance and calculated the principal components using PCA.

As a preliminary analysis, we calculated all the principle components of the Mario domain to see which dimensions PCA weighed the highest during learning. The *jump*, *ground* and *current direction* features are heavily represented in the first few principal components. This is intuitive, as this state changes frequently throughout a game of Mario. These features are also fundamental skills required to play a game of Mario.

PCA also associated features related to enemies within close proximity to Mario entirely in the high-dimensional principal components. These features are only important in very specific scenarios where Mario needs to quickly react to many nearby enemies. Additionally, if only one enemy is nearby, this fact is also represented in the *closest enemy X* and *closest enemy Y* features.

This analysis suggests our approach will work in the Mario Benchmarking Domain because some broadly useful skills can be initially learned, and then refined with additional state dimensions. We will show that we can learn the initial skills quickly. The skills represented in the higher principal components are better for strict optimization in specific scenarios. Analyzing the principal components of the demonstrations is a sanity check as well as validation of the richness of the demonstration.

When using our approach in the Mario domain, results were as hoped (see Figure 4). By projecting the state down to a low-dimensional manifold (less than 4), the learning algorithm converged quickly to bad policies. However, when using a manifold of 4 dimensions or greater, we converged quickly to a much better policy. These are promising results although these well-performing dimensions may still converge to a suboptimal policy after more than 5000 episodes.

Since learning was poor in the first two manifolds, it shows us that the *jump* and *ground* features are not informative enough alone to learn an effective Mario policy. Yet, when projecting onto a 3 or 4 dimensional manifold, we see large increase in policy performance. In these manifolds, the features *jump*, *ground*, *current direction*, *shoot*, *closest enemy Y*, and *obstacles* are all represented. It




Figure 3: The emphasis of each feature relative to the principal component.

is intuitive that these features are important for having basic skill in Mario. The remaining features are important, but only for fine tuning policies.

6. Discussion

The strengths of our approach is its simplicity and generality combined with fast convergence. Once we sample the domain, PCA can be run on these samples. We can use the transforms given by PCA in the learning cycle whenever we compute a new state. This is a simple matrix operation and adds little additional computational cost. This preliminary work shows that learning can be efficiently performed in these low dimensional spaces at an increased convergence rate.

However, our algorithm does make some strong assumptions. First, our algorithm assumes that all state variables are important. It is possible for a state variable to be unnecessary to learning, but have high variance. Second, we stop learning once we converge in the lower-dimensional space. This ensures that given infinite time, the higher dimensional space will converge to better performance. We discuss our solution to these problems in the following section.

7. Future Work

The preliminary results discussed above opens up many potential directions for future research. The next step we will take will be to perform additional analysis on the current results. We want to see if the quality and quantity of demonstration data affects learning. Further experiments include varying the amount of training data given to PCA and using random demonstration data.

A major focus of future work is to improve upon the converged performance of the low-dimensional state representations. Work by Grzes and Kudenko has shown that mixed resolution function approximation works well in complex domains. They learned with a less expressive function approximation to provide early guidance during learning. Simultaneously, they learned using a more expressive function approximation to learn a high quality policy. We propose to leverage a similar idea: we can learn until convergence in a n dimensional space and then transfer to an $n + 1$

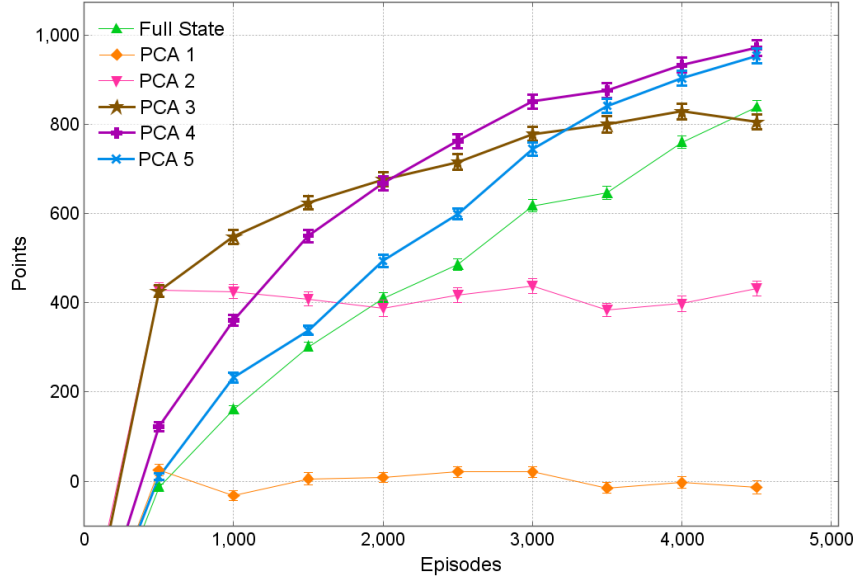


Figure 4: Results with varying manifolds. Lines in bold are experiments that performed better or equal to learning with the full state. Dimensions greater than 5 performed similarly to the full state space, and were not included for clarity. The (small) error bars represent error in the mean over 100 statistical runs.

dimensional space. As shown in our results, the lower dimensional spaces should be faster to learn. We hypothesize that this learning technique will converge faster, and with higher performance, than learning entirely in the full dimensional space.

References

- Mazda Ahmadi, Matthew E. Taylor, and Peter Stone. IFSA: Incremental Feature-Set Augmentation for Reinforcement Learning Tasks. In *Proceedings of the the Sixth International Joint Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, pages 1120–1127, May 2007.
- Brenna D. Argall, Sonia Chernova, Manuela Veloso, and Brett Browning. A survey of robot learning from demonstration. *Robot. Auton. Syst.*, 57(5):469–483, May 2009.
- Tim Brys, Anna Harutyunyan, Peter Vrancx, Matthew E Taylor, Daniel Kudenko, and Ann Nowé. Multi-objectivization of reinforcement learning problems by reward shaping. In *International Joint Conference on Neural Networks (IJCNN)*, pages 2315–2322. IEEE, 2014.
- Tim Brys, Anna Harutyunyan, Halit Bener Suay, Sonia Chernova, Matthew E. Taylor, and Ann Nowé. Reinforcement learning from demonstration through shaping. In *Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI)*, 2015.
- M. Grzes and D. Kudenko. Reinforcement learning with reward shaping and mixed resolution function approximation. *International Journal of Agent Technologies and Systems (IJATS)*, 1(2): 6–54.

- Sergey Karakovskiy and Julian Togelius. The Mario AI benchmark and competitions. *Computational Intelligence and AI in Games, IEEE Transactions on*, 4(1):55–67, 2012.
- Geoffrey Lee, Min Luo, Fabio Zambetta, and Xiaodong Li. Learning a super mario controller from examples of human play. In *Evolutionary Computation (CEC), 2014 IEEE Congress on*, pages 1–8. IEEE, 2014.
- Yizheng Liao, Kun Yi, and Zhe Yang. Cs229 final report reinforcement learning to play mario. Technical report, Stanford University, USA, 2012.
- Bo Liu and Sridhar Mahadevan. Compressive reinforcement learning with oblique random projections, 2011.
- Sinno Jialin Pan and Qiang Yang. A survey on transfer learning. *Knowledge and Data Engineering, IEEE Transactions on*, 22(10):1345–1359, Oct 2010.
- Mykola Pechenizkiy, Seppo Puuronen, and Alexey Tsymbal. Feature extraction for classification in knowledge discovery systems. In *Knowledge-Based Intelligent Information and Engineering Systems*, volume 2773, pages 526–532. Springer, 2003.
- Jonathon Shlens. A tutorial on principal component analysis. In *Systems Neurobiology Laboratory, Salk Institute for Biological Studies*, 2005.
- Christian D. Swinehart and L. F. Abbott. Dimensional reduction for reward-based learning, 2005.
- Matthew E. Taylor and Peter Stone. Representation Transfer for Reinforcement Learning. In *AAAI 2007 Fall Symposium on Computational Approaches to Representation Change during Learning and Development*, November 2007.
- Matthew E. Taylor and Peter Stone. Transfer Learning for Reinforcement Learning Domains: A Survey. *Journal of Machine Learning Research*, 10(1):1633–1685, 2009.
- Matthew E. Taylor, Peter Stone, and Yaxin Liu. Transfer Learning via Inter-Task Mappings for Temporal Difference Learning. *Journal of Machine Learning Research*, 8(1):2125–2167, 2007.
- M.A. Turk and A.P. Pentland. Face recognition using eigenfaces. In *Computer Vision and Pattern Recognition, 1991. Proceedings CVPR '91., IEEE Computer Society Conference on*, pages 586–591, Jun 1991.